

## Exercise 1

**Repeat exercise 9 of Part 3 (text classification with MLPs), now using a bi-directional**

**stacked RNN (with GRU or LSTM cells) with self-attention, all implemented in Keras/TensorFlow or PyTorch.**

### Steps:

1. Implement a bi-directional stacked RNN using GRU or LSTM cells.
2. Implement self-attention (either an MLP or single dense layer) to obtain attention scores.
3. Tune hyperparameters (e.g., number of stacked RNNs, number of hidden layers in self-attention MLP, dropout probability) on the development subset of your dataset.
4. Monitor the performance of the RNN on the development subset during training to decide on the number of epochs to use.
5. Optionally, add an extra RNN layer to produce word embeddings from characters, concatenating the resulting character-based word embeddings with pre-trained word embeddings (e.g., Word2Vec).
6. Include experimental results from:
  - Baseline majority classifier.
  - Best probabilistic classifier from exercise 15 of Part 2.
  - MLP classifier from exercise 9 of Part 3 (treated as a baseline).

### Report:

- Curves showing the loss on training and development data as a function of epochs.
- Precision, recall, F1, precision-recall AUC scores for each class and classifier:
  - Separate for the training, development, and test subsets.
- Macro-averaged precision, recall, F1, precision-recall AUC scores:
  - Averaging the corresponding scores over the classes, separately for the training, development, and test subsets.
- Description of the methods and datasets used:

- Include statistics like average document length, number of training/dev/test documents, and vocabulary size.
- Describe preprocessing steps performed.
- Optionally, try ensemble methods (e.g., majority voting of the best checkpoints, temporal averaging of the weights of the best checkpoints).

## Import Libraries

## Assert whether PyTorch can use an available GPU card

## Creating a Dataset

We will use the `Dataset` class from `PyTorch` to handle the text data. We will pad the text sequences with 0 to a pre-defined length (the average number of tokens in the training split).

```
In [ ]: df_merge = pd.concat([df_fake, df_true], axis = 0 )  
df_merge.head(10)
```

Out[ ]:

	title	text	subject	date	label
0	Donald Trump Sends Out Embarrassing New Year'...	Donald Trump just couldn t wish all Americans ...	News	December 31, 2017	1
1	Drunk Bragging Trump Staffer Started Russian ...	House Intelligence Committee Chairman Devin Nu...	News	December 31, 2017	1
2	Sheriff David Clarke Becomes An Internet Joke...	On Friday, it was revealed that former Milwauk...	News	December 30, 2017	1
3	Trump Is So Obsessed He Even Has Obama's Name...	On Christmas day, Donald Trump announced that ...	News	December 29, 2017	1
4	Pope Francis Just Called Out Donald Trump Dur...	Pope Francis used his annual Christmas Day mes...	News	December 25, 2017	1
5	Racist Alabama Cops Brutalize Black Boy While...	The number of cases of cops brutalizing and ki...	News	December 25, 2017	1
6	Fresh Off The Golf Course, Trump Lashes Out A...	Donald Trump spent a good portion of his day a...	News	December 23, 2017	1
7	Trump Said Some INSANELY Racist Stuff Inside ...	In the wake of yet another court decision that...	News	December 23, 2017	1
8	Former CIA Director Slams Trump Over UN Bully...	Many people have raised the alarm regarding th...	News	December 22, 2017	1
9	WATCH: Brand-New Pro-Trump Ad Features So Muc...	Just when you might have thought we d get a br...	News	December 21, 2017	1

```
In [ ]: #Removing columns which are not required
df = df_merge.drop(["title", "subject","date"], axis = 1)

df.reset_index(inplace = True)
df.drop(["index"], axis = 1, inplace = True)
df.head()
```

Out[ ]:		text	label
0	Donald Trump just couldn t wish all Americans ...		1
1	House Intelligence Committee Chairman Devin Nu...		1
2	On Friday, it was revealed that former Milwauk...		1
3	On Christmas day, Donald Trump announced that ...		1
4	Pope Francis used his annual Christmas Day mes...		1

```
In [ ]:
```

Training set class distribution: [17133 18785]  
Validation set class distribution: [2142 2348]  
Test set class distribution: [2142 2348]

In [ ]:


In [ ]:

Total documents: 44898  
Fake news: 23481, True news: 21417  
Number of training documents: 35918  
Number of validation documents: 4490  
Number of test documents: 4490  
Average document length in training set (in words): 400  
Average document length in validation set (in words): 396  
Average document length in test set (in words): 402  
Vocabulary size (unique words): 397481

## Define the model

We will create a model class and parameterize our neural network with several choices

In [ ]: `plot_loss_curves(results)`




In [ ]:


	Model	Accuracy_train	Accuracy_val	Accuracy_test	Weighted_F1_train \	
0	RNN	0.81	0.79	0.80	0.81	
1	GRU	0.95	0.94	0.94	0.95	
2	LSTM	0.91	0.90	0.90	0.91	
3	LSTM_GMP	1.00	1.00	1.00	1.00	
4	RandomInit	1.00	1.00	1.00	1.00	

	Weighted_F1_val	Weighted_F1_test	PR_AUC_train	PR_AUC_val	PR_AUC_test
0	0.79	0.80	0.894212	0.878705	0.877910
1	0.94	0.94	0.969749	0.963957	0.962466
2	0.90	0.90	0.946631	0.932909	0.933250
3	1.00	1.00	0.999973	0.999362	0.998946
4	1.00	1.00	0.999277	0.998853	0.998003



## Baseline Dummy and Wide MLP



## Model Performance Comparison: Detailed Analysis

The following table summarizes the comparison of five different models using key metrics like accuracy, weighted F1 score, and PR AUC.

Model	Accuracy_train	Accuracy_val	Accuracy_test	Weighted_F1_train	Weighted_F1_val
RNN	0.81	0.79	0.80	0.81	0.79
GRU	0.95	0.94	0.94	0.95	0.94
LSTM	0.91	0.90	0.90	0.91	0.90
LSTM_GMP	1.00	1.00	1.00	1.00	1.00
RandomInit	1.00	1.00	1.00	1.00	1.00

### 1. Training Accuracy:

- **RNN:** The RNN model has the lowest training accuracy at 0.81, indicating that it struggles to capture the underlying patterns in the data compared to other models.
- **GRU:** The GRU model achieves a high training accuracy of 0.95, showing its ability to generalize well from the training data.
- **LSTM:** The LSTM model performs well with 0.91 accuracy, but it is still behind GRU.

- **LSTM\_GMP**: This model achieves perfect training accuracy (1.00), suggesting it has captured the patterns in the training data without overfitting.
- **RandomInit**: Like LSTM\_GMP, the RandomInit model also achieves perfect training accuracy (1.00), which may suggest that its training dynamics are ideal for the task.

## 2. Validation Accuracy:

- **RNN**: The RNN model lags behind with a validation accuracy of 0.79, highlighting that it struggles to generalize beyond the training data.
- **GRU**: The GRU model performs very well on validation data with an accuracy of 0.94, indicating robust generalization.
- **LSTM**: With a validation accuracy of 0.90, LSTM also shows strong generalization, though it still lags behind GRU.
- **LSTM\_GMP** and **RandomInit**: Both of these models achieve perfect validation accuracy (1.00), indicating excellent generalization and minimal overfitting. These models are the most well-suited to handle unseen data.

## 3. Test Accuracy:

- **RNN**: The test accuracy of RNN at 0.80 suggests that its performance on unseen data is not as good as the other models, which is typical for simpler models or ones with insufficient capacity.
- **GRU**: GRU achieves 0.94 test accuracy, performing well on unseen data. This indicates that GRU can maintain good performance across both the training and test sets.
- **LSTM**: With a test accuracy of 0.90, LSTM remains a strong performer, but it falls slightly short of GRU in this case.
- **LSTM\_GMP** and **RandomInit**: Both models excel here with a perfect test accuracy of 1.00, confirming that they are highly effective and have achieved optimal performance.

## 4. Weighted F1 Scores:

- **RNN**: The weighted F1 score for RNN on training, validation, and test sets is around 0.80, which indicates decent performance but room for improvement.
- **GRU**: GRU performs significantly better with a weighted F1 score of 0.95 on training, 0.94 on validation, and 0.94 on test sets, reflecting its strong ability to balance precision and recall.
- **LSTM**: LSTM shows solid F1 scores with values around 0.91, indicating that it balances precision and recall well but still doesn't match GRU in this regard.
- **LSTM\_GMP** and **RandomInit**: Both models reach perfect weighted F1 scores of 1.00 across all datasets, showcasing exceptional ability to classify both classes without bias.

## 5. PR AUC Scores:

- **RNN:** The PR AUC scores for RNN (around 0.88-0.89) suggest that while it has a good overall performance, it isn't as strong as the other models in terms of precision-recall tradeoff.
- **GRU:** The GRU model achieves outstanding PR AUC scores of around 0.96 on both the training, validation, and test sets, indicating excellent performance in distinguishing between positive and negative samples.
- **LSTM:** LSTM's PR AUC scores (around 0.93-0.94) are strong, though they fall short of GRU's performance.
- **LSTM\_GMP** and **RandomInit:** These models achieve near-perfect PR AUC scores, indicating excellent precision-recall tradeoff and almost perfect classification performance.

## 6. Overall Performance Summary:

- **Best Performers:** **LSTM\_GMP** and **RandomInit** stand out as the top performers. They achieve perfect accuracy, weighted F1 scores, and near-perfect PR AUC scores. This suggests that they are the best models for the task, with high generalization and minimal overfitting.
- **Strong Contenders:** **GRU** and **LSTM** are also strong models, with high performance across all metrics. GRU seems to slightly outperform LSTM, but both fall short of LSTM\_GMP and RandomInit in terms of perfect scores.
- **RNN:** The RNN model lags behind, particularly in its ability to generalize, as seen from its lower accuracy and F1 scores. It may benefit from further optimization or the addition of more complex features.

## 7. Conclusion:

- If high accuracy and balanced performance across datasets (train, validation, and test) are the goal, **LSTM\_GMP** and **RandomInit** are the best choices.
- **GRU** and **LSTM** provide good results but can be further fine-tuned to match the performance of the more optimized models.
- **RNN** may require further enhancement or may not be suitable for tasks requiring high generalization across different datasets.

This analysis suggests that **LSTM\_GMP** and **RandomInit** offer the best overall performance, but **GRU** and **LSTM** are still competitive options.

## Exercise 2

Repeat exercise 10 of Part 3, now using a bi-directional stacked RNN (with GRU or LSTM cells) implemented in Keras/TensorFlow or PyTorch.

Tune the hyper-parameters (e.g., number of stacked RNNs, dropout probability) on the development subset. Monitor the performance of the RNN on the development subset during training to decide how many epochs to use.

You may optionally add an extra RNN layer to produce word embeddings from characters, concatenating each resulting character-based word embedding with the corresponding pre-trained word embedding (e.g., obtained with Word2Vec).

Include experimental results of a baseline that tags each word with the most frequent tag it had in the training data; for words that were not encountered in the training data, the baseline should return the most frequent tag (over all words) of the training data.

Also include experimental results of your best method of exercise 10 of Part 3, now treated as an additional baseline.

Include in your report:

- Curves showing the loss on training and development data as a function of epochs.
- Precision, recall, F1, precision-recall AUC scores for each class (tag) and classifier, separately for the training, development, and test subsets, as in exercise 10 of Part 3.

- Macro-averaged precision, recall, F1, precision-recall AUC scores for each classifier, separately for the training, development, and test subsets, as in exercise 10 of Part 3.

- A short description of the methods and datasets you used, including statistics about the datasets (e.g., average document length, number of training/dev/test documents, vocabulary size) and a description of the preprocessing steps that you performed.

You may optionally wish to try ensembles (e.g., majority voting of the best checkpoints, temporal averaging of the weights of the best checkpoints).

## Imports

## Data Download

## Data Parsing and Preprocessing


## RNNs Architecture, Training and Evaluation

```
In [ ]: print("Training Shallow BiGRU RNN:")
        shallow_bigrun_model = train_and_evaluate_rnn(ShallowPOS_BiGRU, input_dim=2*word2ve
```



### Training Shallow BiGRU RNN:

Epoch 1/20, Train Loss: 0.6887225906759813, Dev Loss: 0.5903967435945544  
Epoch 2/20, Train Loss: 0.5520552526641731, Dev Loss: 0.5700697781597462  
Epoch 3/20, Train Loss: 0.5220200925596045, Dev Loss: 0.5456684415875223  
Epoch 4/20, Train Loss: 0.4996648524336199, Dev Loss: 0.5343029464695388  
Epoch 5/20, Train Loss: 0.4838684841695945, Dev Loss: 0.5288984306474078  
Epoch 6/20, Train Loss: 0.47105320412354973, Dev Loss: 0.52468763710114  
Epoch 7/20, Train Loss: 0.4585741071319271, Dev Loss: 0.5197153730798784  
Epoch 8/20, Train Loss: 0.4489992992875383, Dev Loss: 0.5155826621411139  
Epoch 9/20, Train Loss: 0.4402568833217936, Dev Loss: 0.5146788086807519  
Epoch 10/20, Train Loss: 0.43218631883180014, Dev Loss: 0.5097653019174299  
Epoch 11/20, Train Loss: 0.41262031580339326, Dev Loss: 0.5080063198891499  
Epoch 12/20, Train Loss: 0.4087294050393979, Dev Loss: 0.5078501275607518  
Epoch 13/20, Train Loss: 0.40768477494625344, Dev Loss: 0.506912693866811  
Epoch 14/20, Train Loss: 0.40545917346088733, Dev Loss: 0.5071164839025727  
Epoch 15/20, Train Loss: 0.4039940044279265, Dev Loss: 0.5080247047102839  
Epoch 16/20, Train Loss: 0.40277171042679344, Dev Loss: 0.5076611610432914  
Epoch 17/20, Train Loss: 0.4023220165183398, Dev Loss: 0.5074030951524439  
Epoch 18/20, Train Loss: 0.4006993975593876, Dev Loss: 0.5076478105738647  
Epoch 19/20, Train Loss: 0.39882236954982153, Dev Loss: 0.5077785165014124  
Epoch 20/20, Train Loss: 0.3982024533171356, Dev Loss: 0.5075912019587997



Classification Report for ShallowPOS_BiGRU:					
	precision	recall	f1-score	support	
ADJ	0.91	0.90	0.91	1794	
ADP	0.84	0.76	0.80	2030	
ADV	0.93	0.86	0.89	1183	
AUX	0.97	0.96	0.97	1543	
CCONJ	0.76	0.32	0.45	736	
DET	0.88	0.79	0.83	1896	
INTJ	0.92	0.74	0.82	121	
NOUN	0.91	0.90	0.90	4123	
NUM	0.65	0.48	0.56	542	
PART	0.68	0.59	0.63	649	
PRON	0.97	0.97	0.97	2166	
PROPN	0.91	0.81	0.85	2076	
PUNCT	0.59	0.90	0.71	3096	
SCONJ	0.82	0.66	0.73	384	
SYM	0.96	0.61	0.74	109	
VERB	0.93	0.92	0.92	2606	
X	0.20	0.02	0.04	42	
_	0.94	0.77	0.85	354	
accuracy			0.84	25450	
macro avg	0.82	0.72	0.75	25450	
weighted avg	0.86	0.84	0.84	25450	

#### AUC Scores for Each Class:


ADJ: 0.9904  
 ADP: 0.9809  
 ADV: 0.9944  
 AUX: 0.9989  
 CCONJ: 0.9569  
 DET: 0.9889  
 INTJ: 0.9881  
 NOUN: 0.9886  
 NUM: 0.9780  
 PART: 0.9820  
 PRON: 0.9994  
 PROPN: 0.9822  
 PUNCT: 0.9657  
 SCONJ: 0.9912  
 SYM: 0.9733  
 VERB: 0.9950  
 X: 0.9000  
 \_: 0.9878

Macro-Averaged Precision: 0.8197999103769109  
 Macro-Averaged Recall: 0.7205777278839265  
 Macro-Averaged F1: 0.7544068783965447  
 Macro-Averaged Precision-Recall AUC: 0.9801045208141689

```
In [ ]: print("Training Somewhat Deep BiGRU RNN:")
        somewhat_deep_bigrun_model = train_and_evaluate_rnn(DeepPOS_BiGRU, input_dim=2*word
```

Training Somewhat Deep BiGRU RNN:

Epoch 1/20, Train Loss: 0.7022546846293019, Dev Loss: 0.5867635689881212  
Epoch 2/20, Train Loss: 0.5667467373823074, Dev Loss: 0.5607323512472305  
Epoch 3/20, Train Loss: 0.5355007804888793, Dev Loss: 0.5470747769924632  
Epoch 4/20, Train Loss: 0.5163764904924496, Dev Loss: 0.5313968856009027  
Epoch 5/20, Train Loss: 0.4987494789519525, Dev Loss: 0.5240710062117206  
Epoch 6/20, Train Loss: 0.4859537860313045, Dev Loss: 0.5268476989334986  
Epoch 7/20, Train Loss: 0.47321452441555695, Dev Loss: 0.5183539290475965  
Epoch 8/20, Train Loss: 0.46294248138149996, Dev Loss: 0.5155142432540879  
Epoch 9/20, Train Loss: 0.45372238668249004, Dev Loss: 0.5186465924843809  
Epoch 10/20, Train Loss: 0.4451829866587269, Dev Loss: 0.5136722904772388  
Epoch 11/20, Train Loss: 0.4223506854107308, Dev Loss: 0.5125919201021505  
Epoch 12/20, Train Loss: 0.4175819735985578, Dev Loss: 0.5130881210018817  
Epoch 13/20, Train Loss: 0.4148119417116405, Dev Loss: 0.5125126642242709  
Epoch 14/20, Train Loss: 0.4130115674381789, Dev Loss: 0.5124499859442389  
Epoch 15/20, Train Loss: 0.4112976084637598, Dev Loss: 0.5140063311895332  
Epoch 16/20, Train Loss: 0.41111166755898343, Dev Loss: 0.513076528869476  
Epoch 17/20, Train Loss: 0.4085400805171807, Dev Loss: 0.5147684721420881  
Epoch 18/20, Train Loss: 0.4067330705513919, Dev Loss: 0.5148656620716391  
Epoch 19/20, Train Loss: 0.40543398701478256, Dev Loss: 0.5143077934072131  
Epoch 20/20, Train Loss: 0.40476068490546674, Dev Loss: 0.5151277575875285



# Classification Report for DeepPOS\_BiGRU:

	precision	recall	f1-score	support
ADJ	0.91	0.91	0.91	1794
ADP	0.84	0.76	0.80	2030
ADV	0.94	0.85	0.89	1183
AUX	0.97	0.96	0.97	1543
CCONJ	0.83	0.31	0.45	736
DET	0.89	0.78	0.83	1896
INTJ	0.93	0.69	0.80	121
NOUN	0.91	0.90	0.90	4123
NUM	0.66	0.49	0.56	542
PART	0.68	0.61	0.65	649
PRON	0.96	0.97	0.97	2166
PROPN	0.91	0.81	0.86	2076
PUNCT	0.58	0.91	0.71	3096
SCONJ	0.80	0.70	0.74	384
SYM	0.96	0.61	0.74	109
VERB	0.93	0.92	0.92	2606
X	0.00	0.00	0.00	42
_	0.96	0.78	0.86	354
accuracy			0.84	25450
macro avg	0.81	0.72	0.75	25450
weighted avg	0.86	0.84	0.84	25450

## AUC Scores for Each Class:

ADJ: 0.9901  
 ADP: 0.9810  
 ADV: 0.9948  
 AUX: 0.9988  
 CCONJ: 0.9573  
 DET: 0.9891  
 INTJ: 0.9802  
 NOUN: 0.9887  
 NUM: 0.9760  
 PART: 0.9829  
 PRON: 0.9994  
 PROPN: 0.9807  
 PUNCT: 0.9663  
 SCONJ: 0.9905  
 SYM: 0.9739  
 VERB: 0.9949  
 X: 0.9072  
 \_: 0.9875

Macro-Averaged Precision: 0.8144968015657622

Macro-Averaged Recall: 0.719536159674343


Macro-Averaged F1: 0.7530539398151068

Macro-Averaged Precision-Recall AUC: 0.9799530944194375

```
In [ ]: print("Training Very Deep BiGRU RNN:")
        very_deep_bigrun_model = train_and_evaluate_rnn(VeryDeepPOS_BiGRU, input_dim=2*word
```

### Training Very Deep BiGRU RNN:

Epoch 1/20, Train Loss: 0.7415308893886593, Dev Loss: 0.6143699082216822  
Epoch 2/20, Train Loss: 0.5946611025962585, Dev Loss: 0.5762313990888739  
Epoch 3/20, Train Loss: 0.5588858117699771, Dev Loss: 0.5607933800919611  
Epoch 4/20, Train Loss: 0.5382228776389173, Dev Loss: 0.5455909664544246  
Epoch 5/20, Train Loss: 0.5208595192831854, Dev Loss: 0.5391640282588495  
Epoch 6/20, Train Loss: 0.5071301893863213, Dev Loss: 0.5371348662558654  
Epoch 7/20, Train Loss: 0.4958442723432612, Dev Loss: 0.525490907088557  
Epoch 8/20, Train Loss: 0.4843710587372045, Dev Loss: 0.5296513514411181  
Epoch 9/20, Train Loss: 0.4747367534254936, Dev Loss: 0.5240084873255632  
Epoch 10/20, Train Loss: 0.46566733328592197, Dev Loss: 0.5202993790159249  
Epoch 11/20, Train Loss: 0.4445159603381982, Dev Loss: 0.513571013782855  
Epoch 12/20, Train Loss: 0.4370920961473075, Dev Loss: 0.5127045885662088  
Epoch 13/20, Train Loss: 0.4349677289771106, Dev Loss: 0.5127411887311099  
Epoch 14/20, Train Loss: 0.432260421029372, Dev Loss: 0.5131513121582213  
Epoch 15/20, Train Loss: 0.42951581396802185, Dev Loss: 0.5139819231621903  
Epoch 16/20, Train Loss: 0.4280960612146298, Dev Loss: 0.5127779939270258  
Epoch 17/20, Train Loss: 0.42704468451721056, Dev Loss: 0.5134377005629074  
Epoch 18/20, Train Loss: 0.425876667353752, Dev Loss: 0.514354779494735  
Epoch 19/20, Train Loss: 0.42580834123144334, Dev Loss: 0.5128783027181649  
Epoch 20/20, Train Loss: 0.4208277455491369, Dev Loss: 0.5134728133379666



Classification Report for VeryDeepPOS\_BiGRU:

	precision	recall	f1-score	support
ADJ	0.91	0.91	0.91	1794
ADP	0.84	0.75	0.79	2030
ADV	0.94	0.86	0.90	1183
AUX	0.96	0.97	0.96	1543
CCONJ	0.71	0.34	0.46	736
DET	0.89	0.79	0.83	1896
INTJ	0.93	0.76	0.84	121
NOUN	0.92	0.89	0.90	4123
NUM	0.66	0.50	0.57	542
PART	0.65	0.63	0.64	649
PRON	0.96	0.98	0.97	2166
PROPN	0.89	0.81	0.85	2076
PUNCT	0.59	0.89	0.71	3096
SCONJ	0.80	0.69	0.74	384
SYM	0.93	0.61	0.73	109
VERB	0.92	0.92	0.92	2606
X	0.14	0.02	0.04	42
_	0.95	0.78	0.86	354
accuracy			0.84	25450
macro avg	0.81	0.73	0.76	25450
weighted avg	0.85	0.84	0.84	25450

AUC Scores for Each Class:

ADJ: 0.9900  
ADP: 0.9806  
ADV: 0.9944  
AUX: 0.9988  
CCONJ: 0.9574  
DET: 0.9889  
INTJ: 0.9820  
NOUN: 0.9883  
NUM: 0.9756  
PART: 0.9828  
PRON: 0.9994  
PROPN: 0.9802  
PUNCT: 0.9663  
SCONJ: 0.9884  
SYM: 0.9711  
VERB: 0.9948  
X: 0.8996  
\_: 0.9865

Macro-Averaged Precision: 0.8111026383944473  
Macro-Averaged Recall: 0.7271724473564507  
Macro-Averaged F1: 0.7571378802757028  
Macro-Averaged Precision-Recall AUC: 0.9791739557097381

## Choice of Architectures for Shallow, Somewhat Deep, and Deep RNNs

The design of the three RNN models (shallow, somewhat deep, and deep) was driven by the goal of maintaining a similar structure to the original MLP models while utilizing recurrent layers (GRU or LSTM) to capture temporal dependencies. The architectures were chosen based on the depth and complexity of the original MLP models, with the following considerations:

---

## 1. Shallow Bidirectional RNN (ShallowPOS\_BiGRU)

- **Depth:** The original shallow MLP contained only a single hidden layer. To maintain a similar structure in the RNN, we opted for **one bidirectional RNN layer**.
  - **Bidirectional Nature:** The bidirectional RNN (using GRU or LSTM) processes the sequence in both directions (forward and backward), capturing context from both past and future words. This enriches the representation without increasing the depth of the model.
  - **Output Layer:** Since the task is POS tagging, we use the final hidden state of the sequence for classification. We take the output of the last timestep from the bidirectional RNN as the representation of the input sequence, which is passed through a final linear layer for POS tagging.
  - **Regularization:** Dropout is included after the RNN layer to prevent overfitting, keeping the model simple and regularized.
- 

## 2. Somewhat Deep Bidirectional RNN (DeepPOS\_BiGRU)

- **Depth:** The original somewhat deep MLP contained two hidden layers, and we retained the concept of depth by adding **two bidirectional RNN layers**.
  - **Bidirectional Nature:** Just like the shallow model, the bidirectional RNN layers help capture both past and future context. With two layers, the model can learn richer representations of the input sequence.
  - **Regularization:** Dropout is applied after each RNN layer, just like the original MLP's design, to prevent overfitting while maintaining simplicity.
  - **Output Layer:** Similar to the shallow model, we take the final hidden state from the last timestep after the second RNN layer as the representation for POS tagging.
- 

## 3. Very Deep Bidirectional RNN (VeryDeepPOS\_BiGRU)

- **Depth:** The original very deep MLP had **three hidden layers**. To replicate this depth in the RNN model, we used **three bidirectional RNN layers**.
- **Bidirectional Nature:** With three bidirectional RNN layers, the model learns contextual information from both the past and the future at multiple levels of abstraction.
- **Regularization:** Dropout is applied after each RNN layer to reduce the risk of overfitting and ensure that the model generalizes well.

- **Output Layer:** As in the shallower models, the last timestep of the final RNN layer is used for classification, which provides the most relevant representation of the entire sequence.

Each RNN architecture was designed to match the depth and structure of the original MLP models while incorporating the ability of recurrent layers to capture sequential dependencies. The bidirectional nature of the RNNs enhances the model's performance by considering both past and future contexts for each input sequence.

Model	Accuracy	Macro avg F1- score	Weighted avg F1- score	Strengths	Weaknesses
ShallowPOS_MLP	0.83	0.74	0.83	High performance on <b>PRON (0.97 precision)</b> , <b>AUX (0.96 recall)</b> , and <b>VERB (0.91 f1-score)</b> .	Struggles with <b>X (f1 = 0.00)</b> and <b>CCONJ (0.45 f1)</b> .
DeepPOS_MLP	0.83	0.74	0.83	Improved recall for <b>AUX (0.94)</b> and <b>PRON (0.97)</b> , maintains strong performance across categories.	Struggles with rare categories like <b>X</b> .
VeryDeepPOS_MLP	0.82	0.74	0.82	Strong performance on <b>AUX</b> and <b>PRON</b> , but slightly lower overall performance compared to other MLP models.	Some decrease in performance, especially in <b>CCONJ (0.43 f1)</b> and <b>NUM (0.56 f1)</b> .
ShallowPOS_BiGRU	0.84	0.75	0.84	High performance on <b>PRON (0.97 precision)</b> , <b>AUX (0.97 recall)</b> , and <b>VERB (0.92 f1-score)</b> .	Struggles with <b>X (f1 = 0.04)</b> and <b>CCONJ (0.45 f1)</b> .
DeepPOS_BiGRU	0.84	0.75	0.84	High performance on <b>AUX (0.97)</b> and <b>PRON (0.97)</b> .	Struggles with <b>X</b> and <b>CCONJ (0.45 f1)</b> .



Model	Accuracy	Macro avg F1- score	Weighted avg F1- score	Strengths	Weaknesses
VeryDeepPOS_BiGRU	0.84	0.76	0.84	Strong performance on <b>AUX (0.97)</b> , <b>PRON (0.98)</b> , and <b>VERB (0.92)</b> with a good balance across categories.	Some degradation in performance for <b>CCONJ (0.46 f1)</b> and <b>NUM (0.57 f1)</b> , but more consistent than previous models.

### Comparison:

- **Accuracy:** Both the MLP and BiGRU models generally perform similarly in terms of accuracy (ranging from **0.82 to 0.84**).
- **F1 Scores:** The BiGRU models tend to outperform the MLP models slightly in terms of both macro and weighted average F1 scores, with values like **0.75 (macro avg)** and **0.84 (weighted avg)** compared to the MLP's **0.74 (macro avg)** and **0.83 (weighted avg)**.
- **Category Performance:** Both MLP and RNN models show consistent strength in categories like **PRON** and **AUX**, but also share similar weaknesses in rare categories like **X** and **CCONJ**.
- **Deepness & Performance Trade-off:** Moving from shallow to deep MLP models generally resulted in small drops in accuracy and macro F1, suggesting diminishing returns with deeper layers. However, the BiGRU models show slight improvements in their macro and weighted F1 scores despite being deep.

### Baseline Tagger

In [ ]:

### Baseline Tagger Classification Report:

	precision	recall	f1-score	support
ADJ	0.91	0.83	0.87	1794
ADP	0.87	0.88	0.88	2030
ADV	0.94	0.79	0.86	1183
AUX	0.93	0.89	0.91	1543
CCONJ	0.99	1.00	0.99	736
DET	0.96	0.97	0.96	1896
INTJ	0.97	0.69	0.80	121
NOUN	0.67	0.93	0.78	4123
NUM	0.91	0.61	0.73	542
PART	0.69	0.99	0.81	649
PRON	0.96	0.93	0.95	2166
PROPN	0.91	0.51	0.66	2076
PUNCT	0.99	0.99	0.99	3096
SCONJ	0.62	0.60	0.61	384
SYM	0.81	0.83	0.82	109
VERB	0.89	0.82	0.85	2606
X	1.00	0.00	0.00	42
_	0.97	0.81	0.89	354
accuracy			0.86	25450
macro avg	0.89	0.78	0.80	25450
weighted avg	0.88	0.86	0.86	25450

### AUC Scores for Each Class:

ADJ: 0.9111  
ADP: 0.9350  
ADV: 0.8924  
AUX: 0.9422  
CCONJ: 0.9985  
DET: 0.9823  
INTJ: 0.8429  
NOUN: 0.9219  
NUM: 0.8029  
PART: 0.9902  
PRON: 0.9629  
PROPN: 0.7539  
PUNCT: 0.9927  
SCONJ: 0.7992  
SYM: 0.9170  
VERB: 0.9017  
X: 0.5000  
\_: 0.9066

Macro-Averaged Precision: 0.8891

Macro-Averaged Recall: 0.7814

Macro-Averaged F1: 0.7974

Macro-Averaged Precision-Recall AUC: 0.8863

## Dataset Statistics



Training Data Statistics:

Total number of words: 207230

Vocabulary size: 20201

Average word length: 4.08 characters

Number of unique POS tags: 18

Most frequent POS tag: NOUN (occurred 34755 times)

POS tag distribution:

ADJ: 13187 occurrences

ADP: 17745 occurrences

ADV: 10117 occurrences

AUX: 12818 occurrences

CCONJ: 6687 occurrences

DET: 16299 occurrences

INTJ: 695 occurrences

NOUN: 34755 occurrences

NUM: 4127 occurrences

PART: 5748 occurrences

PRON: 18677 occurrences

PROPN: 12618 occurrences

PUNCT: 23596 occurrences

SCONJ: 3822 occurrences

SYM: 722 occurrences

VERB: 22604 occurrences

X: 399 occurrences

\_: 2614 occurrences

Development Data Statistics:

Total number of words: 25512

Vocabulary size: 5638

Average word length: 4.14 characters

Number of unique POS tags: 18

Most frequent POS tag: NOUN (occurred 4212 times)

POS tag distribution:

ADJ: 1873 occurrences

ADP: 2039 occurrences

ADV: 1224 occurrences

AUX: 1567 occurrences

CCONJ: 779 occurrences

DET: 1900 occurrences

INTJ: 115 occurrences

NOUN: 4212 occurrences

NUM: 383 occurrences

PART: 647 occurrences

PRON: 2225 occurrences

PROPN: 1865 occurrences

PUNCT: 3075 occurrences

SCONJ: 397 occurrences

SYM: 83 occurrences

VERB: 2710 occurrences

X: 59 occurrences

\_: 359 occurrences

Test Data Statistics:

Total number of words: 25450

Vocabulary size: 5750  
Average word length: 4.13 characters  
Number of unique POS tags: 18  
Most frequent POS tag: NOUN (occurred 4123 times)

POS tag distribution:  
ADJ: 1794 occurrences  
ADP: 2030 occurrences  
ADV: 1183 occurrences  
AUX: 1543 occurrences  
CCONJ: 736 occurrences  
DET: 1896 occurrences  
INTJ: 121 occurrences  
NOUN: 4123 occurrences  
NUM: 542 occurrences  
PART: 649 occurrences  
PRON: 2166 occurrences  
PROPN: 2076 occurrences  
PUNCT: 3096 occurrences  
SCONJ: 384 occurrences  
SYM: 109 occurrences  
VERB: 2606 occurrences  
X: 42 occurrences  
\_: 354 occurrences

## Methods and Datasets

We developed a **Part-of-Speech (POS) tagger** using a variety of **Recurrent Neural Network (RNN)** architectures, using either **Bidirectional GRU** or **LSTM** models. In all cases, **Word2Vec embeddings** were used as input features. The model was evaluated on the **English Universal Dependencies Treebank** (UD\_English-EWT), which contains labeled data for training, development, and testing.

The models were evaluated on several performance metrics such as **accuracy**, **precision**, **recall**, **F1-score**, and **AUC scores**.

### Datasets:

- **Training:** en\_ewt-ud-train.conllu
- **Development:** en\_ewt-ud-dev.conllu
- **Test:** en\_ewt-ud-test.conllu

The data was parsed and preprocessed to extract **words** and their corresponding **POS tags**.

## Preprocessing Steps

### 1. Tokenization:

- We parsed the **conllu** files to extract each word and its associated POS tag.

### 2. Word2Vec Embeddings:

- Pre-trained **Word2Vec** embeddings were used to represent words as vectors. This allows us to capture semantic relationships between words.

### 3. **Model:**

- We trained various models, including:
  - **Bidirectional RNN with GRU/LSTM cells:** These architectures were developed to capture sequential dependencies in the data. The models had variety in their depth. Bidirectional GRU/LSTM cells were used to capture both past and future context in the sentence. Dropout was applied as a regularization technique.
  - **Comparison of Performance:** We compared different model architectures—**Shallow RNN**, **Deep RNN**, and **Very Deep RNN** with the previous best MLP performer.

### 4. **Evaluation:**

- The performance of the models was evaluated on **accuracy**, **precision**, **recall**, **F1-score**, and **AUC scores**, with metrics computed separately for each POS tag and averaged across tags.